


Computational Experiments on Task-Based Parallel Applications

Salão de Iniciação Científica UFRGS 2019

Henrique Corrêa Pereira da Silva
Lucas Mello Schnorr (advisor)



UNIVERSIDADE FEDERAL DO RIO GRANDE DO SUL



Context

In the last few years, we saw the continuation of a decades long trend: ever-growing **parallelism** in search of performance



Context

In the last few years, we saw the continuation of a decades long trend: ever-growing **parallelism** in search of performance

Also we saw that. . .



Context

In the last few years, we saw the continuation of a decades long trend: ever-growing **parallelism** in search of performance

Also we saw that. . .

- the hardware naturally specialized over time



Context

In the last few years, we saw the continuation of a decades long trend: ever-growing **parallelism** in search of performance

Also we saw that. . .

- the hardware naturally specialized over time
- buying highly performant systems became very expensive



Context

In the last few years, we saw the continuation of a decades long trend: ever-growing **parallelism** in search of performance

Also we saw that. . .

- the hardware naturally specialized over time
- buying highly performant systems became very expensive
- more common systems configurations became hard to extract their full potential



Context

In the last few years, we saw the continuation of a decades long trend: ever-growing **parallelism** in search of performance

Also we saw that. . .

- the hardware naturally specialized over time
- buying highly performant systems became very expensive
- more common systems configurations became hard to extract their full potential

Applications running on them can no longer rely on homogeneous hardware if they seek *high performance*



Reality

Thus, a solution is needed in face of this demand. . .



Reality

Thus, a solution is needed in face of this demand. . .

Some sort of **middle layer** capable of partitioning applications workloads into these hybrid systems!



Agenda

1 Introduction

2 Method

- Proof-of-concept implementation
- Experimental validation

3 Results

- Visualizations
- Conclusions

Agenda

1 Introduction

2 Method

- Proof-of-concept implementation
- Experimental validation

3 Results

- Visualizations
- Conclusions

Most popular

There are several popular parallel programming models



Most popular

There are several popular parallel programming models

Examples:



Most popular

There are several popular parallel programming models

Examples:

- Message passing



Most popular

There are several popular parallel programming models

Examples:

- Message passing
- Shared memory



Most popular

There are several popular parallel programming models

Examples:

- Message passing
- Shared memory
- Hybrid model



Most popular

There are several popular parallel programming models

Examples:

- Message passing
- Shared memory
- Hybrid model

In any of those models, their implementation is a manual, complex, time-consuming and, therefore, error-prone process



Implementations

Most popular implementations seek to alleviate the burden of programmers



Implementations

Most popular implementations seek to alleviate the burden of programmers

Of those implementations the most popular are:



Implementations

Most popular implementations seek to alleviate the burden of programmers

Of those implementations the most popular are:

- OpenMP
- OpenMPI
- OpenACC
- etc..



Implementations

Most popular implementations seek to alleviate the burden of programmers

Of those implementations the most popular are:

- OpenMP
- OpenMPI
- OpenACC
- etc..

Attention!

If only utilizing the previously cited APIs, the domain decomposition is normally fixed to the number of resources



Implementations

Most popular implementations seek to alleviate the burden of programmers

Of those implementations the most popular are:

- OpenMP
- OpenMPI
- OpenACC
- etc..

Attention!

If only utilizing the previously cited APIs, the domain decomposition is normally fixed to the number of resources **you'll be victim to dynamic load imbalances**



StarPU

One of the efforts into creating a middleware capable of overcoming said limitations is called **StarPU**




One of the efforts into creating a middleware capable of overcoming said limitations is called **StarPU**

Its approach is by defining the problem by tasks, and those tasks into a *Directed Acyclic Graph* (or *DAG* for short)




One of the efforts into creating a middleware capable of overcoming said limitations is called **StarPU**


Its approach is by defining the problem by tasks, and those tasks into a *Directed Acyclic Graph* (or *DAG* for short)



Task-based



Task-based



Agenda

1 Introduction

2 Method

- Proof-of-concept implementation
- Experimental validation

3 Results

- Visualizations
- Conclusions

Intent

Our objectives were to



Intent

Our objectives were to

- 1 Use the task-based approach



Intent

Our objectives were to

- 1 Use the task-based approach
- 2 Learn the StarPU API



Intent

Our objectives were to

- 1 Use the task-based approach
- 2 Learn the StarPU API
- 3 Analyze its behavior while executing the application




Object of study

A simple vector accumulation



Object of study

A simple vector accumulation



Example kernel

```
void reduc_sum(void** buffers, void* cl_arg)
{
    ullint* vec_input = (ullint*)STARPU_VECTOR_GET_PTR(buffers[0]);
    ullint* output = (ullint*)STARPU_VARIABLE_GET_PTR(buffers[1]);
    uint nx_input = STARPU_VECTOR_GET_NX(buffers[0]);

    double t0 = get_time();

    // do the job
    for (uint i = 0; i < nx_input; i++)
        *output += vec_input[i];

    double t1 = get_time();

    V_PRINTF("SUM = %d\n",
            "Task finished work with elapsed time %f\n",
            *output, t1 - t0);
}
```



Agenda

1 Introduction

2 Method

- Proof-of-concept implementation
- Experimental validation

3 Results

- Visualizations
- Conclusions

Basis of comparison

Other simple implementations of vector accumulations



Basis of comparison

Other simple implementations of vector accumulations

- 1 Naive
- 2 C++ STL
- 3 OpenMP



Basis of comparison

Other simple implementations of vector accumulations

- 1 Naive
- 2 C++ STL
- 3 OpenMP

Not intended as a goal to beat, but as a basis



Methodological approach

A full factorial, randomly ordered experiment design



Methodological approach

A full factorial, randomly ordered experiment design

Parameters:

- *Vector size*: $7 \cdot 10^7$, $3 \cdot 10^8$ and $1.1 \cdot 10^9$
- *Number of blocks*: 7000, 25000 and 82000
- *Reduction factor*: 2, 10 and 1000



Agenda

1 Introduction


2 Method

- Proof-of-concept implementation
- Experimental validation

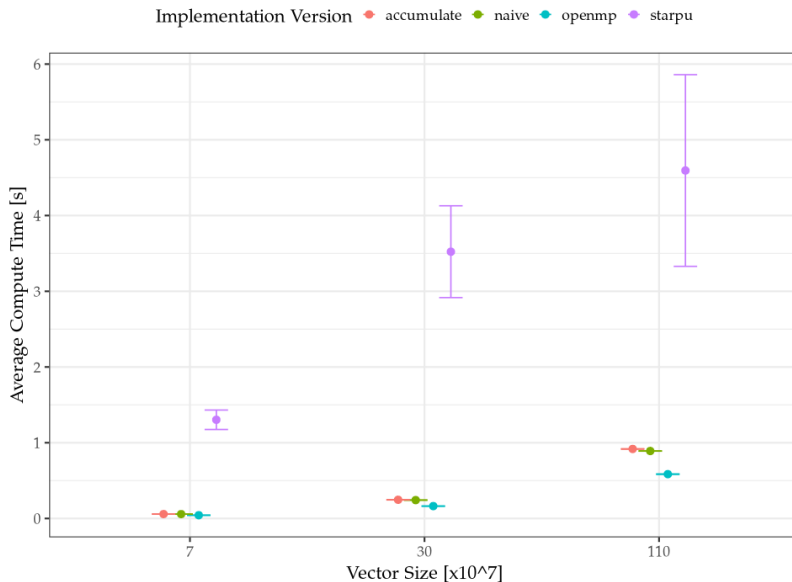
3 Results

- Visualizations
- Conclusions

Number of Blocks ◆ 7000 ◆ 25000 ◆ 82000



Combined graph



Agenda

1 Introduction

2 Method

- Proof-of-concept implementation
- Experimental validation

3 Results

- Visualizations
- Conclusions

Feasibility

Even with a simple implementation, we have shown that **StarPU** is a very capable API



Feasibility

Even with a simple implementation, we have shown that **StarPU** is a very capable API

Furthermore, we can aggregate the other APIs into our computation kernel



Feasibility

Even with a simple implementation, we have shown that **StarPU** is a very capable API

Furthermore, we can aggregate the other APIs into our computation kernel

OpenMP therefore, utilizing *parallel tasks*

OpenMPI distribute the execution graph across a whole cluster



Questions?

Salão de Iniciação Científica UFRGS 2019

Henrique Corrêa Pereira da Silva
Lucas Mello Schnorr (advisor)